

TAP CSR Research Agent

V1 → V2 Change Report — what changed, and why

Fundraising · The Apprentice Project

Executive summary

V1 proved the concept: an automated researcher that reads a company's public CSR record across six sources and scores its fit with TAP's mission, with every claim carrying a source URL and verbatim excerpt. V2 keeps that zero-hallucination architecture and rebuilds the judgment on top of it. The score is now calibrated against TAP's real partners and real non-fits, the team's manual research methodology (8 criteria, 0–5) is built into the tool and is the single source of the verdict, evidence must sit in CSR/education context to count, unsourced assertions are banned from every output, and the app is deployed as a serverless Flask application on Vercel.

1. Where V1 stood

Six research sources (India CSR page, MCA CSR-2 via CIN, National CSR Portal, annual report PDF, funded partners, LinkedIn decision-makers). A six-dimension 0–100 keyword score: Focus 40 · Adjacency 20 · Geography 10 · Maturity 10 · Budget 10 · Source quality 10. Two modes (Prospect Screening, Deep Research), DOCX and HTML reports, Streamlit UI on Streamlit Cloud. Roughly 2,800 lines of Python.

2. Why a V2 — the five failures that drove it

Real partners under-scored. Capgemini — an actual TAP partner — scored 77/100 because its CSR pages say 'digital inclusion' and 'digital academy', not TAP's keyword list. Keyword matching missed meaning.

Non-fits over-scored. Apollo Tyres (trucker-healthcare CSR) reached 60 and Diageo (vocational 'Learning for Life') reached 75 — stray keywords earned education points with no sense of proportion or intent.

Evidence out of context. The keyword 'maker' matched a smartphone-patent story as STEM evidence; 'education' matched a website navigation bar; a LinkedIn profile's NatWest experience defined Ericsson's delivery model.

Three verdicts from one run. The DOCX led with 92/100 'Priority Match', the XLSX with 'Worth Hunting, 3.0/5', the HTML carried both. Two fundraisers acting on different files would do opposite things.

Deployment ceiling. Streamlit requires a long-lived server and cannot run on Vercel, the organisation's chosen platform.

3. Scoring engine changes

Dimension	V1	V2	Why
Focus alignment	40	40	Adjacency-aware uplift; softened depth penalty; higher points for education-supporter keywords
Adjacency boost	20	18	Rebalanced slightly toward partner similarity
Partner similarity	—	15	NEW — 8 traits researched from TAP's six real partners

Geography fit	10	9	Rebalanced
CSR maturity	10	8	Rebalanced
Budget size	10	5	A big budget must not mask weak alignment
Source quality	10	5	Same

On top of the weights, four new judgment layers: (a) an AI semantic layer (Groq / Llama-3.3-70B) that reads the fetched evidence against TAP's mission and the researched profile of TAP's actual partners — it can only raise the focus dimension, never lower it; (b) an education-dominance dampener that measures what share of the CSR portfolio is actually about education and scales the education-driven dimensions down when it is a side theme; (c) a rule-based employability detector that caps vocational-dominant portfolios at 50, independent of the LLM; and (d) an indirect-alignment ceiling of 89 so the top tier stays reserved for direct mission matches.

4. Grounded in TAP's real partners

We researched the public CSR records of TAP's actual partners — Tenerity, Capgemini, Bajaj Finserv, NRB Bearings, SVP India and Project Tech4Dev — and distilled eight shared traits (funds NGO-implemented programmes, digital inclusion/skilling, STEM and robotics for government-school students, skilled employee volunteering, WhatsApp/chatbot delivery, children-and-youth soft skills, NGO capacity building). These traits form the new partner-similarity dimension: a prospect can now reach Strong Fit either by direct mission keywords or by looking like the companies that already fund TAP.

5. The manual methodology, built in

The team's manual research template (8 criteria scored 0–5; verdict tiers Priority Hunt / Worth Hunting / Conditional Fit / Low Priority; 7-sheet deep-dive base; strict source hierarchy; gaps marked 'to confirm', never filled) is now derived automatically from the same fetched evidence. Critically, **the scorecard is the single source of the verdict**: the website, DOCX, HTML, XLSX and JSON all lead with the same tier and 0–5 average, and the 0–100 engine score is presented as a diagnostic only. One run produces one verdict.

6. Evidence discipline and trust policy

Context windows. A keyword counts only within ± 300 characters of CSR/education context in the same passage — and never inside navigation/menu-style text.

Names must be names. Decision-makers must match a person-name pattern with a junk-word blocklist; role fields are cut at the next name boundary; the project ledger rejects career-history strings.

Right company only. Delivery-model and programme evidence excludes people-search/LinkedIn profile text entirely.

No assertion without a source. The unsourced 'existing TAP partner' claim is removed everywhere; the internal partner list no longer affects the score and survives only as a JSON annotation labelled as tool configuration. If a claim cannot be traced to fetched text, it does not render.

Honest verification. The verification log now distinguishes VERIFIED (string exists AND sits in supporting context) from CONTEXT UNVERIFIED and CHECK MANUALLY — string existence alone no longer produces 100% pass rates.

7. Calibration results

Test case	V1-era	V2	Why it moved
Capgemini-like (real partner profile)	77	83	Semantic + partner-similarity credit
Apollo Tyres-like (health-only CSR)	60	40	Education-dominance dampener
Diageo-like (vocational skilling)	75	50	Dampener + employability cap
Healthcare control (must stay low)	13	13	Controls hold — no inflation

Numbers from the repeatable calibration suite (synthetic evidence profiles run through both engines, LLM disabled for determinism).

8. Platform

The UI was rewritten from Streamlit to a Flask application (api/index.py + webui.py) deployed serverless on Vercel with a 300-second Fluid-compute budget. Because serverless functions share no memory between requests, Deep Research generates all four report files in-request and embeds them in the results page as downloads — stateless and immune to the V1 'download resets the page' bug. The research pipeline modules were untouched.

9. What's next

Cross-funder summary scorecard and Loom-script draft (steps 2–3 of the methodology); a standing benchmark of ~12 known companies run on every calibration change; continued tightening of extraction quality. The verify step remains the job: the engine drafts, a person checks every figure against the source before anything ships.